

5-Minute Demo Video — Full Script

Includes the required 60–120 second code walkthrough segment. Timed to 5:00.

BEFORE YOU HIT RECORD

- Have the app open, logged in, on the landing page.
- Have a second window ready with `sql/07_ensemble/070_final_match_scores.sql` and `agents/02_price_optimization_agent.sql` open for the code walkthrough.
- Backup matched pair: **Canon CB-2LV Battery Charger -- Abt ID 36053, Buy ID 90141687.**
- **Pre-record the Market Intelligence Agent's answer separately** -- its response has taken 40s to 6+ minutes in testing.
- Do one full silent run-through with a stopwatch before recording.

0:00–0:30

Hook

SHOW

- Dashboard landing page, hero banner visible.

SAY

“This is our submission for the Snowflake Cortex Hackathon: an AI-Powered Product Matching System for the E-commerce and Retail industry. Retailers need to know which of their products a competitor also sells, and at what price. We solve this as an entity-resolution problem — matching products across two independent catalogs using four AI signals combined into one explainable score, not a single similarity guess.”

0:30–0:45

Dashboard First Impression

SHOW

- Scroll to the pipeline row, click one step to show the popover.

SAY

“Here's our dashboard, built entirely on Streamlit in Snowflake, querying real data live. Five pipeline steps, fully interactive — click any one to see exactly what it does.”

0:45–2:15

Code Walkthrough (90s) [REQUIRED SEGMENT]

SHOW

- Switch to the GitHub repo / editor file tree.
- Show folder structure: `sql/`, `agents/`, `semantic_model/`, `mcp/`, `streamlit/`.
- Open `sql/07_ensemble/070_final_match_scores.sql`, scroll to the `ensemble-combine + GREEDY_1TO1_RESOLVE` procedure.
- Open `agents/02_price_optimization_agent.sql`, show the `tool_resources` block.

SAY

“Let's look at the actual code. The repository is organized as a numbered SQL pipeline: raw data loading, blocking, embeddings, attribute extraction, LLM adjudication, and ensemble scoring — each stage in its own file, in the order it runs.”

“Here's the core of the differentiator: the ensemble scoring script. It combines four independent signals — semantic similarity, structured attribute similarity, and a cost-gated LLM verdict for ambiguous cases — into one weighted confidence score. Below that, a Python stored procedure resolves conflicts, so no product is ever claimed by two matches at once.”

“Each of our three Cortex Agents is defined declaratively in its own file. This one, the Price Optimization Agent, wires a deterministic pricing rule engine as a callable tool — the AI only phrases the justification, it never invents the price itself.”

“The Streamlit dashboard, the semantic view, and the MCP server all read from these same underlying tables — one pipeline output, reused everywhere, with no duplicated logic.”

2:15–2:50

Accuracy Proof

SHOW

- Matching Accuracy tab, click the “What is Precision?” popover, open a false-positive example.

SAY

“Here's the proof, measured live against the real ground-truth mapping: 97.9% precision, 86.6% recall, F1 of 0.919. When this system says two products are the same, it's right 98% of the time.”

“And when it's wrong, you can see exactly why — the actual embedding similarity, the actual attribute match. That's the difference between a similarity score and a real explanation.”

2:50–3:10

Product Explorer

SHOW

- Search “sony”, then the side-by-side comparison with the Canon CB-2LV pair.

SAY

“Universal search across both catalogs, plus side-by-side comparison for any two products. Confidence score, rationale, done.”

3:10–3:40

Product Matching Agent — Cortex Agent #1 (live)

SHOW

- PRODUCT_MATCHING_AGENT chat. Ask: “Why are Canon CB-2LV Battery Charger products a match?”

SAY

“This is Cortex Agent number one: the Product Matching Agent, live in Snowflake Intelligence. Watch what it cites: 0.977 confidence, a perfect attribute match on the manufacturer part number, and 0.954 semantic similarity on the listing text. It's reasoning over the actual signal values, live, right now.”

3:40–4:05

Price Optimization Agent — Cortex Agent #2 (live)

SHOW

- PRICE_OPTIMIZATION_AGENT chat. Ask for a pricing recommendation on the same pair.

SAY

“Cortex Agent number two: Price Optimization. The pricing engine is a deterministic rule engine, deliberately — the AI only phrases the justification, never invents the price. Hold, undercut, or raise, with the reasoning stated plainly.”

4:05–4:25

Market Intelligence Agent — Cortex Agent #3 (pre-recorded)

SHOW

- Cut to the pre-recorded clip showing its category-trend answer.

SAY

“Cortex Agent number three: Market Intelligence, using AI_AGG to synthesize a narrative across an entire category in one call — because not every insight needs to wait on an LLM to render on the dashboard itself.”

4:25–4:40

Ask Anything

SHOW

- Type a question, expand “See the query used”.

SAY

“Ask anything about this data in plain English. It writes its own SQL, validates it, and answers from the real result — every query is auditable, right here.”

4:40–4:50

MCP Integration

SHOW

- Show an MCP client connected to ABT_BUY_MCP_SERVER (or skip and redistribute time if untested).

SAY

“This entire system is also exposed through a managed MCP server, so any MCP-compatible client can query it directly — the same architecture powering the dashboard and the agents.”

4:50–5:00

Close

SHOW

- Landing page or a plain closing slide.

SAY

“Every match explainable. Every price justified. Every trend narrated. Built entirely on Snowflake Cortex. Thank you for watching.”

***Total: 5:00**, including the required 60–120s code walkthrough. If running long, the Code Walkthrough and Accuracy Proof segments carry the most evaluation weight — trim Product Explorer or Ask Anything first instead.*